



Transformers and Pre-trained Language Models





Danqi Chen

Princeton Language and Intelligence
Princeton University

July 15th, 2024

(Many slides are adapted from Princeton COS484 and Stanford CS224N course materials)

Introduction



Adithya
Bhaskar

Howard Chen
(w/ Karthik
Narasimhan)

Dan Friedman

Tianyu Gao


Lucy He
(w/ Peter
Henderson)



Alexander
Wettig



Mengzhou Xia



Howard Yen



Zexuan Zhong

Our research focuses on ***training, adapting and understanding*** large language models

See more at <https://www.cs.princeton.edu/~danqic/>

Lecture plan

Part I. **Transformers**

Focus: innovations and key designs in neural **architectures**



Part II. **Pre-trained language models**

Focus: training **objectives & data**, downstream **adaptations**

Lecture plan

- **Fundamentals (70%)** - I will walk through the most important ideas in NLP and LLMs in the past 5+ years (Transformers, pre-training, in-context learning, RLHF, ...)
- **How do these ideas evolve and lead to state-of-the-art models? (15%)**
 - I will highlight recent improvements and developments
- **Cutting-edge research topics (15%)** - What research topics do we study in 2024? I will briefly discuss some of the works from my research group too

Part I.

Transformers

Transformers

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com



(Vaswani et al., 2017)

What is attention?

NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau


Jacobs University Bremen, Germany

KyungHyun Cho Yoshua Bengio*

Université de Montréal

(Bahdanau et al., 2015)

- Attention is a technique to address the “bottleneck” issue in the seq2seq model, originally designed for machine translation



What is attention?

NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau

Jacobs University Bremen, Germany


KyungHyun Cho Yoshua Bengio*

Université de Montréal


(Bahdanau et al., 2015)

- Attention is a technique to address the “bottleneck” issue in the seq2seq model, originally designed for machine translation
- **Key idea**  : At each time step during decoding, focus on only a particular part of source sentence
 - ▶ This depends on **decoder's** current hidden state h_t^{dec}
 - ▶ Usually implemented as a probability distribution over the hidden states of the encoder (h_i^{enc})

Attention for seq2seq models



Attention for seq2seq models




Attention learns the notion of **alignment**

“Which source words are more relevant to the current target word?”

Attention for seq2seq models

$h_1^{enc}, \dots, h_n^{enc}$ and h_t^{dec} are hidden states from encoder and decoder RNNs




- Encoder hidden states: $h_1^{enc}, \dots, h_n^{enc}$ (n: # of words in source sentence)
- Decoder hidden state at time t : h_t^{dec}
- Attention scores:
$$e^t = [g(h_1^{enc}, h_t^{dec}), \dots, g(h_n^{enc}, h_t^{dec})] \in \mathbb{R}^n$$
- Attention distribution:
$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^n$$
- Weighted sum of encoder hidden states:
$$o_t = \sum_{i=1}^n \alpha_i^t h_i^{enc} \in \mathbb{R}^h$$

Combine o_t and h_t^{dec} to predict next word


Attention as a soft, averaging lookup table

We can think of **attention** as performing fuzzy lookup a in **key-value store**

Lookup table: a table of **keys** that map to **values**. The **query** matches one of the keys, returning its value.




Attention: The **query** matches to all **keys** softly to a weight between 0 and 1. The keys' **values** are multiplied by the weights and summed.




(In the case of NMT, key = value)

Transformer encoder-decoder




(Vaswani et al., 2017)

- Transformer encoder + Transformer decoder:
 - a replacement for seq2seq + attention based on RNNs
- First designed and experimented on NMT



Transformers (both encoders and decoders) have become the default neural architectures in modeling languages!

Transformer encoder-decoder



- Transformer encoder = a stack of **encoder layers**
- Transformer decoder = a stack of **decoder layers**

Transformer encoder: BERT, RoBERTa, ELECTRA

Transformer decoder: GPT-n, ChatGPT, Gemini, Claude, LLaMA, Mistral, ...


Transformer encoder-decoder: T5, BART

- Key innovation: **self-attention, multi-head**
- Transformers don't have any recurrence structures!

$$\mathbf{h}_t = f(\mathbf{h}_{t-1}, \mathbf{x}_t) \in \mathbb{R}^h$$

(Vaswani et al., 2017)

Transformers: roadmap



- Self-attention and multi-head attention
- Feedforward layers
- Positional encoding
- Residual connections + layer normalization
- Transformer encoder vs Transformer decoder

The Annotated Transformer

Attention is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com


Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

<https://nlp.seas.harvard.edu/annotated-transformer/>

Transformers: roadmap



- Self-attention and multi-head attention ←
- Feedforward layers
- Positional encoding
- Residual connections + layer normalization
- Transformer encoder vs Transformer decoder

- Advanced techniques: SwiGLU, rotary embeddings, pre-normalization, grouped query attention
- Architecture exploration beyond Transformers

General form of attention

- A more general form: use a set of **keys** and **values** $(\mathbf{k}_1, \mathbf{v}_1), \dots, (\mathbf{k}_n, \mathbf{v}_n)$, $\mathbf{k}_i \in \mathbb{R}^{d_k}, \mathbf{v}_i \in \mathbb{R}^{d_v}$, **keys** are used to compute the attention scores and **values** are used to compute the output vector
- Attention always involves the following steps:
 - Computing the **attention scores** $\mathbf{e} = g(\mathbf{q}, \mathbf{k}_i) \in \mathbb{R}^n$
 - Taking softmax to get **attention distribution** α :



$$\alpha = \text{softmax}(\mathbf{e}) \in \mathbb{R}^n$$

- Using attention distribution to take **weighted sum** of values:

$$\mathbf{o} = \sum_{i=1}^n \alpha_i \mathbf{v}_i \in \mathbb{R}^{d_v}$$

Self-attention

- In NMT, **query** = decoder's hidden state, **keys** = **values** = encoder's hidden states
- Self-attention = attention from the sequence to **itself**
- Self-attention: let's use each word in a sequence as the **query**, and all other words in the sequence as **keys** and **values**.



Self-attention

Step #1: Transform each input vector into three vectors: **query**, **key**, and **value** vectors

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q \in \mathbb{R}^{d_q}$$


$$\mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K \in \mathbb{R}^{d_k}$$

$$\mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V \in \mathbb{R}^{d_v}$$

$$\mathbf{W}^Q \in \mathbb{R}^{d_{in} \times d_q}$$

$$\mathbf{W}^K \in \mathbb{R}^{d_{in} \times d_k}$$

$$\mathbf{W}^V \in \mathbb{R}^{d_{in} \times d_v}$$



Note that we use row vectors here;
It is also common to write

$$\mathbf{q}_i = \mathbf{W}^Q \mathbf{x}_i \in \mathbb{R}^{d_q}$$

for \mathbf{x}_i = a column vector

Self-attention

Step #2: Compute pairwise similarities between keys and queries; normalize with softmax

For each \mathbf{q}_i , compute attention scores and attention distribution:


$$\alpha_{i,j} = \text{softmax}\left(\frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}\right)$$

aka. “scaled dot product”

It must be $d_q = d_k$ in this case

Q. Why scaled dot product?


To avoid the dot product to become too large
for larger d_k ; scaling the dot product by $\frac{1}{\sqrt{d_k}}$
is easier for optimization



Self-attention

Step #3: Compute output for each input as weighted sum of values

$$\mathbf{h}_i = \sum_{j=1}^n \alpha_{i,j} \mathbf{v}_j \in \mathbb{R}^{d_v}$$






Self-attention

What would be the output vector for the word “Thinking” approximately?

- (A) $0.5\mathbf{v}_1 + 0.5\mathbf{v}_2$
- (B) $0.54\mathbf{v}_1 + 0.46\mathbf{v}_2$
- (C) $0.88\mathbf{v}_1 + 0.12\mathbf{v}_2$
- (D) $0.12\mathbf{v}_1 + 0.88\mathbf{v}_2$

(C) is correct.



Self-attention: matrix notations

$X \in \mathbb{R}^{n \times d_{in}}$ (n = input length)

$$Q = XW^Q, K = XW^K, V = XW^V$$

where $W^Q \in \mathbb{R}^{d_{in} \times d_q}, W^K \in \mathbb{R}^{d_{in} \times d_k}, W^V \in \mathbb{R}^{d_{in} \times d_v}$

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Diagram illustrating the components of the self-attention formula:

- Q is a matrix of size $n \times d_q$.
- K is a matrix of size $d_k \times n$.
- V is a matrix of size $n \times d_v$.

A red arrow points from the text "Q: What is this softmax operation?" to the softmax term in the equation.

$$\text{softmax}\left(\frac{\begin{matrix} Q & K^T \\ \times & \end{matrix}}{\sqrt{d_k}}\right) V$$

Diagram illustrating the softmax operation:


- Q is a matrix of size $n \times d_q$.
- K^T is a matrix of size $d_k \times n$.
- \times indicates matrix multiplication between Q and K^T .
- $\sqrt{d_k}$ is the denominator of the softmax function.
- V is a matrix of size $n \times d_v$.

The result of the softmax operation is labeled H , which is a matrix of size $n \times d_v$.

Multi-head attention

“The Beast with Many Heads”

- It is better to use multiple attention functions instead of one!
 - Each attention function (“head”) can focus on different positions.
- It gives the attention layer multiple “representation subspaces”




Multi-head attention

“The Beast with Many Heads”

Finally, we just concatenate all the heads and apply an output projection matrix.

$$\begin{aligned}\text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \\ \text{head}_i &= \text{Attention}(XW_i^Q, XW_i^K, XW_i^V)\end{aligned}$$



- In practice, we use a **reduced** dimension for each head.

$$W_i^Q \in \mathbb{R}^{d_{in} \times d_q}, W_i^K \in \mathbb{R}^{d_{in} \times d_k}, W_i^V \in \mathbb{R}^{d_{in} \times d_v}$$

$$d_q = d_k = d_v = d/m \quad d = \text{hidden size, } m = \# \text{ of heads}$$

$$W^O \in \mathbb{R}^{d \times d_{out}}$$

- The total computational cost is similar to that of single-head attention with full dimensionality.

Multi-head attention

“The Beast with Many Heads”



$$\begin{aligned}\text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_n)W^O \\ \text{head}_i &= \text{Attention}(XW_i^Q, XW_i^K, XW_i^V)\end{aligned}$$

- We can think of multi-head attention (MHA) layer as an abstraction layer that maps a sequence of input vectors $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^{d_{in}}$ to a sequence of n vectors: $\mathbf{h}_1, \dots, \mathbf{h}_n \in \mathbb{R}^{d_{out}}$
If we stack multiple layers, usually $d_{in} = d_{out} = d$
- The same abstraction as RNNs - used as a drop-in replacement for an RNN layer


$$\mathbf{h}_t = f(\mathbf{W}\mathbf{h}_{t-1} + \mathbf{U}\mathbf{x}_t + \mathbf{b}) \in \mathbb{R}^{d_{out}}$$

Much easier to parallelize, more expensive to scale up to longer sequences!

What does multi-head attention learn?



Transformers: roadmap



- Self-attention and multi-head self-attention
- Feedforward layers ←
- Positional encoding
- Residual connections + layer normalization
- Transformer encoder vs Transformer decoder
- Advanced techniques: SwiGLU, rotary embeddings, pre-normalization, grouped query attention
- Architecture exploration beyond Transformers

Adding nonlinearities: Feed-forward layers


- There are no element-wise nonlinearities in self-attention; stacking more self-attention layers just re-averages value vectors
- Simple fix: add a feed-forward network to post-process each output vector

$$\text{FFN}(\mathbf{x}_i) = \text{ReLU}(\mathbf{x}_i \mathbf{W}_1 + \mathbf{b}_1) \mathbf{W}_2 + \mathbf{b}_2$$


$$\mathbf{W}_1 \in \mathbb{R}^{d \times d_{ff}}, \mathbf{b}_1 \in \mathbb{R}^{d_{ff}}$$

$$\mathbf{W}_2 \in \mathbb{R}^{d_{ff} \times d}, \mathbf{b}_2 \in \mathbb{R}^d$$

Usually, $d_{ff} = 4d$



Transformers: roadmap



- Self-attention and multi-head attention
- Feedforward layers
- Positional encoding
- Residual connections + layer normalization
- Transformer encoder vs Transformer decoder
- Advanced techniques: SwiGLU, rotary embeddings, pre-normalization, grouped query attention
- Architecture exploration beyond Transformers


Modeling order information: positional encoding

- Unlike RNNs, self-attention doesn't build in order information, we need to encode the order of the sentence in our keys, queries, and values
- Solution: Add **positional embeddings** to the input embeddings: $\mathbf{p}_i \in \mathbb{R}^d$ for $i = 1, 2, \dots, n$

$$\mathbf{x}_i \leftarrow \mathbf{x}_i + \mathbf{p}_i$$

- **Sinusoidal positional embeddings:** sine and cosine functions of different frequencies:


$$p_i = \begin{pmatrix} \sin(i/10000^{2*1/d}) \\ \cos(i/10000^{2*1/d}) \\ \vdots \\ \sin(i/10000^{2*\frac{d}{2}/d}) \\ \cos(i/10000^{2*\frac{d}{2}/d}) \end{pmatrix}$$



- **Pros:** Periodicity + can extrapolate to longer sequences
- **Cons:** Not learnable


Modeling order information: positional encoding

- **Absolute positional embeddings:** let all \mathbf{p}_i be learnable parameters
 - $P \in \mathbb{R}^{d \times L}$ for $L = \text{max sequence length}$
 - **Pros:** each position gets to be learned to fit the data
 - **Cons:** can't extrapolate to indices outside of max sequence length L
 - **Examples:** BERT, GPT-1, GPT-2, GPT-3, OPT



(Devlin et al., NAACL 2019)

Transformers: roadmap



- Self-attention and multi-head attention
- Feedforward layers
- Positional encoding
- Residual connections + layer normalization ←
- Transformer encoder vs Transformer decoder
- Advanced techniques: SwiGLU, rotary embeddings, pre-normalization, grouped query attention
- Architecture exploration beyond Transformers

How to make Transformers work for **deep** NNs?

Add & Norm: $\text{LayerNorm}(x + \text{Sublayer}(x))$

Residual connections (He et al., CVPR 2016)

Instead of $X^{(i)} = \text{Layer}(X^{(i-1)})$ (i represents the layer)



We let $X^{(i)} = X^{(i-1)} + \text{Layer}(X^{(i-1)})$, so we only need to learn “the residual” from the previous layer



This prevents the network from "forgetting" or distorting important information as it is processed by many layers.

How to make Transformers work for **deep** NNs?

Add & Norm: $\text{LayerNorm}(x + \text{Sublayer}(x))$


Layer normalization (Ba et al., 2016)

Problem: Difficult to train the parameters of a given layer because its input from the layer beneath keeps shifting.

Solution: Reduce variation by **normalizing** to zero mean and standard deviation of one **within each layer**.


$$y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta \quad \gamma, \beta \in \mathbb{R}^d \text{ are learnable parameters}$$

Transformers: roadmap



- Self-attention and multi-head attention
- Feedforward layers
- Positional encoding
- Residual connections + layer normalization
- Transformer encoder vs Transformer decoder ←
- Advanced techniques: SwiGLU, rotary embeddings, pre-normalization, grouped query attention
- Architecture exploration beyond Transformers

Let's put things together - Transformer encoder



From the bottom to the top:


- Input embedding
- Positional encoding
- A stack of Transformer encoder layers

Transformer encoder is a stack of N layers, which consists of two sub-layers:

- Multi-head attention layer
- Feed-forward layer

$$\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^{d_{in}} \longrightarrow \mathbf{h}_1, \dots, \mathbf{h}_n \in \mathbb{R}^{d_{out}}$$

Let's put things together - Transformer decoder



From the bottom to the top:


- Output embedding
- Positional encoding
- A stack of Transformer decoder layers
- Linear + softmax

Transformer decoder is a stack of N layers, which consists of **three** sub-layers:

- Masked multi-head attention
- Multi-head cross-attention
- Feed-forward layer
- (w/ Add & Norm between sub-layers)

Masked multi-head self-attention

- Key: You can't see the future text for the decoder!




- Solution: for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$ **How to implement this? Masking!**

Masked multi-head self-attention

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q, \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K, \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$$

$$e_{i,j} = \frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}, \forall j = 1, \dots, n$$

$$\alpha_i = \text{softmax}(\mathbf{e}_i)$$



Efficient implementation: compute attention as we normally do, mask out attention to future words by setting attention scores to $-\infty$

```
dot = torch.bmm(queries, keys.transpose(1, 2))

indices = torch.triu_indices(t, t, offset=1)
dot[:, indices[0], indices[1]] = float('-inf')

dot = F.softmax(dot, dim=2)
```



Masked multi-head self-attention

The following matrix denotes the values of $\frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}$ for $1 \leq i \leq n, 1 \leq j \leq n$ ($n = 4$)


| | | | |
|----|----|----|----|
| 1 | 0 | -1 | -1 |
| 1 | 1 | -1 | 0 |
| 0 | 1 | 1 | -1 |
| -1 | -1 | 2 | 1 |

What should be the value of $\alpha_{2,2}$ in masked attention?


- (A) 0
- (B) 0.5
- (C) $\frac{e}{2e + e^{-1} + 1}$
- (D) 1

The correct answer is (B)

Multi-head cross-attention



Similar as the attention in seq2seq model!



Multi-head cross-attention

Self-attention:

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q, \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K, \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$$

$$e_{i,j} = \frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}, \forall j = 1, \dots, n$$

$$\alpha_i = \text{softmax}(\mathbf{e}_i)$$

$$\mathbf{h}_i = \sum_{j=1}^n \alpha_{i,j} \mathbf{v}_j$$

Cross-attention:

(always from the top layer)

$\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_m$: hidden states from encoder

$\mathbf{x}_1, \dots, \mathbf{x}_n$: hidden states from decoder

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q \quad i = 1, 2, \dots, n$$


$$\mathbf{k}_j = \tilde{\mathbf{x}}_j \mathbf{W}^K, \mathbf{v}_j = \tilde{\mathbf{x}}_j \mathbf{W}^V \quad \forall j = 1, 2, \dots, m$$

$$e_{i,j} = \frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}, \forall j = 1, \dots, m$$

$$\alpha_i = \text{softmax}(\mathbf{e}_i)$$

$$\mathbf{h}_i = \sum_{j=1}^m \alpha_{i,j} \mathbf{v}_j$$

Transformer encoder-decoder



Training Transformer encoder-decoder models

The same as the way that we train seq2seq models!

- Training data: parallel corpus $\{(\mathbf{w}_i^{(s)}, \mathbf{w}_i^{(t)})\}$
- Minimize cross-entropy loss:

$$\sum_{t=1}^T -\log P(y_t | y_1, \dots, y_{t-1}, \mathbf{w}^{(s)})$$

(denote $\mathbf{w}^{(t)} = y_1, \dots, y_T$)

- Back-propagate gradients through both encoder and decoder

36M sentence pairs

French: bonjour le monde .



English: hello world .

Masked self-attention is the key!

This can enable parallelizable operations while NOT looking at the future


Empirical results with Transformers

| Model | BLEU | | Training Cost (FLOPs) | |
|---------------------------------|-------------|--------------|-----------------------|---------------------------------------|
| | EN-DE | EN-FR | EN-DE | EN-FR |
| ByteNet [15] | 23.75 | | | |
| Deep-Att + PosUnk [32] | | 39.2 | | $1.0 \cdot 10^{20}$ |
| GNMT + RL [31] | 24.6 | 39.92 | $2.3 \cdot 10^{19}$ | $1.4 \cdot 10^{20}$ |
| ConvS2S [8] | 25.16 | 40.46 | $9.6 \cdot 10^{18}$ | $1.5 \cdot 10^{20}$ |
| MoE [26] | 26.03 | 40.56 | $2.0 \cdot 10^{19}$ | $1.2 \cdot 10^{20}$ |
| Deep-Att + PosUnk Ensemble [32] | | 40.4 | | $8.0 \cdot 10^{20}$ |
| GNMT + RL Ensemble [31] | 26.30 | 41.16 | $1.8 \cdot 10^{20}$ | $1.1 \cdot 10^{21}$ |
| ConvS2S Ensemble [8] | 26.36 | 41.29 | $7.7 \cdot 10^{19}$ | $1.2 \cdot 10^{21}$ |
| Transformer (base model) | 27.3 | 38.1 | | $3.3 \cdot 10^{18}$ |
| Transformer (big) | 28.4 | 41.0 | | $2.3 \cdot 10^{19}$ |

(Vaswani et al., 2017)

Transformer-based language models

The backbone of large language models (e.g., GPT/ChatGPT, Gemini, LLaMA, ...)




Transformer architecture specifications

| | N | d_{model} | d_{ff} | h | d_k | d_v |
|------|-----|--------------------|-----------------|-----|-------|-------|
| base | 6 | 512 | 2048 | 8 | 64 | 64 |

(Vaswani et al., 2017)

| Model Name | n_{params} | n_{layers} | d_{model} | n_{heads} | d_{head} |
|-----------------------|---------------------|---------------------|--------------------|--------------------|-------------------|
| GPT-3 Small | 125M | 12 | 768 | 12 | 64 |
| GPT-3 Medium | 350M | 24 | 1024 | 16 | 64 |
| GPT-3 Large | 760M | 24 | 1536 | 16 | 96 |
| GPT-3 XL | 1.3B | 24 | 2048 | 24 | 128 |
| GPT-3 2.7B | 2.7B | 32 | 2560 | 32 | 80 |
| GPT-3 6.7B | 6.7B | 32 | 4096 | 32 | 128 |
| GPT-3 13B | 13.0B | 40 | 5140 | 40 | 128 |
| GPT-3 175B or “GPT-3” | 175.0B | 96 | 12288 | 96 | 128 |

(Brown et al., 2020)



Transformers: pros and cons

- **Easier to capture long-range dependencies:** we draw attention between every pair of words!
- **Easier to parallelize:**

$$Q = XW^Q \quad K = XW^K \quad V = XW^V$$

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

- **Are positional embeddings enough to capture positional information?**

Otherwise self-attention is an unordered function of its input

- **Quadratic computation in self-attention**

Can become very slow when the sequence becomes very long

Computational analysis of Transformers

Multi-head attention (MHA)

$$Q = XW^Q, K = XW^K, V = XW^V$$
$$\begin{matrix} n \times d \\ d \times n \\ n \times d \end{matrix}$$

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

$$O(nd^2 + n^2d)$$


Feed-forward layers (FFN)

$$\begin{matrix} d \times d_{ff} \\ d_{ff} \times d \\ n \times d \end{matrix}$$
$$\text{FFN}(\mathbf{x}_i) = \text{ReLU}(\mathbf{x}_i \mathbf{W}_1 + \mathbf{b}_1) \mathbf{W}_2 + \mathbf{b}_2$$
$$O(nd^2)$$

Note: RNNs only require $O(nd^2)$ time: $\mathbf{h}_t = f(\mathbf{W}\mathbf{h}_{t-1} + \mathbf{U}\mathbf{x}_t + \mathbf{b})$
(assuming input dimension = hidden dimension = d)


Computational analysis of Transformers

- For BERT-sized models ($n = 512$, $d = 768$, $d_{ff} = 4d$), 2/3 of parameters are FFNs.



- However, when sequence length becomes longer (e.g., $> 50,000$), the computation will be dominated by self-attention $O(n^2d)$
 - Numerous solutions have been proposed to address this issue
 - Long-context language modeling is still one of the most active research areas today

Transformers: roadmap



- Self-attention and multi-head attention
- Feedforward layers
- Positional encoding
- Residual connections + layer normalization
- Transformer encoder vs Transformer decoder


- Advanced techniques: SwiGLU, rotary embeddings, pre-normalization, grouped query attention
- Architecture exploration beyond Transformers

Major modifications since original
Transformers


SwiGLU activation

$$\text{SwiGLU} = \text{Swish} + \text{GLU}$$

$$\text{SwiGLU}(x, W, V, b, c, \beta) = \text{Swish}_\beta(xW + b) \otimes (xV + c)$$



$$\text{Swish}(x) = x \cdot \text{sigmoid}(\beta x)$$



<https://azizbelaweid.substack.com/p/what-is-swiglu-how-to-implement-it>


SwiGLU activation

| | Score | CoLA | SST-2 | MRPC | MRPC | STSB | STSB | QQP | QQP | MNLI _m | MNLI _{mm} | QNLI | RTE |
|-------------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|-------------------|--------------------|--------------|--------------|
| | Average | MCC | Acc | F1 | Acc | PCC | SCC | F1 | Acc | Acc | Acc | Acc | Acc |
| FFN _{ReLU} | 83.80 | 51.32 | 94.04 | 93.08 | 90.20 | 89.64 | 89.42 | 89.01 | 91.75 | 85.83 | 86.42 | 92.81 | 80.14 |
| FFN _{GELU} | 83.86 | 53.48 | 94.04 | 92.81 | 90.20 | 89.69 | 89.49 | 88.63 | 91.62 | 85.89 | 86.13 | 92.39 | 80.51 |
| FFN _{Swish} | 83.60 | 49.79 | 93.69 | 92.31 | 89.46 | 89.20 | 88.98 | 88.84 | 91.67 | 85.22 | 85.02 | 92.33 | 81.23 |
| FFN _{GLU} | 84.20 | 49.16 | 94.27 | 92.39 | 89.46 | 89.46 | 89.35 | 88.79 | 91.62 | 86.36 | 86.18 | 92.92 | 84.12 |
| FFN _{GEGLU} | 84.12 | 53.65 | 93.92 | 92.68 | 89.71 | 90.26 | 90.13 | 89.11 | 91.85 | 86.15 | 86.17 | 92.81 | 79.42 |
| FFN _{Bilinear} | 83.79 | 51.02 | 94.38 | 92.28 | 89.46 | 90.06 | 89.84 | 88.95 | 91.69 | 86.90 | 87.08 | 92.92 | 81.95 |
| FFN _{SwiGLU} | 84.36 | 51.59 | 93.92 | 92.23 | 88.97 | 90.32 | 90.13 | 89.14 | 91.87 | 86.45 | 86.47 | 92.93 | 83.39 |
| FFN _{ReGLU} | 84.67 | 56.16 | 94.38 | 92.06 | 89.22 | 89.97 | 89.85 | 88.86 | 91.72 | 86.20 | 86.40 | 92.68 | 81.59 |
| [Raffel et al., 2019] | 83.28 | 53.84 | 92.68 | 92.07 | 88.92 | 88.02 | 87.94 | 88.67 | 91.56 | 84.24 | 84.57 | 90.48 | 76.28 |
| ibid. stddev. | 0.235 | 1.111 | 0.569 | 0.729 | 1.019 | 0.374 | 0.418 | 0.108 | 0.070 | 0.291 | 0.231 | 0.361 | 1.393 |

$$\text{SwiGLU}(x, W, V, b, c, \beta) = \text{Swish}_\beta(xW + b) \otimes (xV + c)$$

Notes: there are 3 projection matrices (up_project, down_project, gate_project), d_{ff} is reduced to $4d \times \frac{2}{3}$

Pre-normalization



RMSNorm normalization function

$$\bar{a}_i = \frac{a_i}{\text{RMS}(\mathbf{a})} g_i, \quad \text{where } \text{RMS}(\mathbf{a}) = \sqrt{\frac{1}{n} \sum_{i=1}^n a_i^2}.$$

(Zhang and Senrich, 2019)

Image: (Xiong et al., 2020)

Rotary positional embeddings

- **Relative positional embeddings (T5 uses this!):**

Self-Attention with Relative Position Representations

Peter Shaw
Google
petershaw@google.com


Jakob Uszkoreit
Google Brain
usz@google.com

Ashish Vaswani
Google Brain
avaswani@google.com

$$e_{ij} = \frac{x_i W^Q (x_j W^K + a_{ij}^K)^T}{\sqrt{d_z}}$$

- Instead of focusing on absolute positions, relative positional embeddings concentrate on the **distances between pairs of tokens**
- Incorporating this relative positional information into attention directly

Rotary positional embeddings




Unites both **absolute** and **relative** positional information

$$f_{\{q,k\}}(\mathbf{x}_m, m) = \begin{pmatrix} \cos m\theta & -\sin m\theta \\ \sin m\theta & \cos m\theta \end{pmatrix} \begin{pmatrix} W_{\{q,k\}}^{(11)} & W_{\{q,k\}}^{(12)} \\ W_{\{q,k\}}^{(21)} & W_{\{q,k\}}^{(22)} \end{pmatrix} \begin{pmatrix} x_m^{(1)} \\ x_m^{(2)} \end{pmatrix}$$

(Su et al., 2021) RoFormer: Enhanced Transformer with Rotary Position Embedding


Grouped query attention (GQA)



(Ainslie et al., 2023) GQA: Training generalized multi-query transformer models from multi-head checkpoints.


Architecture exploration beyond Transformers

Efficient Transformers



(Tay et al., 2020): Efficient Transformers: A Survey

Example: Performers




L : sequence length, $m \ll L$


Low-rank decomposition: Decompose A as the product of Q' and K' (random projection of original keys and querys)

Example: Longformer / Big Bird


Sparse attention: only compute attention at particular positions




(a) Full n^2 attention



(b) Sliding window attention




(c) Dilated sliding window




(d) Global+sliding window


(Beltagy et al., 2020): Longformer: The Long-Document Transformer




(a) Random attention



(b) Window attention



(c) Global Attention




(d) BIGBIRD


(Zaheer et al., 2021): Big Bird: Transformers for Longer Sequences

Example: Transformer-XL

Segment-level recurrence with state reuse: hidden representations from previous segment will be cached as extended context (no back-propagation to those!)




Research from my group



(Zhong et al., EMNLP'22) Training Language Models with Memory Augmentation





AutoCompressors



(Chevalier et al., EMNLP'23) Adapting Language Models to Compress Contexts

Research from my group

CEPE




These architectures/techniques are generally applicable to both **long-context modeling** and **retrieval augmentation!**

(Yen et al., ACL'24) Long-Context Language Modeling with Parallel Context Encoding

Remarks on efficient Transformers


- A lot of exploration around 2019-2021: mostly approximation solutions of replacing the full quadratic attention. Few techniques have been adopted in state-of-the-art LLMs (exception: Mistral uses *Sliding Window Attention* to handle longer sequences).



You can still do such approximations to speed up inference though!


- On the other hand, many system-level advancements have been made to scale up Transformers to longer sequences **without approximation** e.g., FlashAttention (Dao et al., 2022)

Vision Transformer (ViT)






(Dosovitskiy et al., 2021): An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale

Mixture of experts (MoEs)



Mixture of experts (MoEs)



$$y = \text{FFN}(x; \sum_i r_i \theta_i)$$

$r_i = \text{softmax}(W_r x)$

Merged FFN

Routing

Part II.

Pre-trained language models

Roadmap

- **The BERT era:** pre-training and fine-tuning
- **The GPT-3 era:** prompting and in-context learning
- **The ChatGPT era:** supervised instruction tuning and RLHF



The BERT era: pre-training and fine-tuning

BERT = Bidirectional Encoder Representations from Transformers

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova
Google AI Language

{jacobdevlin, mingweichang, kentonl, kristout}@google.com

(Devlin et al., 2019)

Input: a sequence of n words

Output: a sequence of n vectors

aka. “**contextualized word embeddings**”

Each word doesn’t have a fixed vector as in (static) word embeddings

LSTMs →



ELMo





← Transformer decoder

GPT-1



BERT

Inside BERT: A Transformer encoder



- 12 or 24 layers of Transformer blocks
- Each block consists of a multi-head self-attention layer and a feedforward layer with residual connections

How is BERT pre-trained?

Two **pre-training** objectives:

- Masked language modeling (MLM)
- Next sentence prediction (NSP)

Masked language modeling

= mask out **15%** of the inputs and
then predict the masked words

store



gallon



the man went to the [MASK] to buy a [MASK] of milk

- Rather than *always* replacing the chosen words with [MASK], the data generator will do the following:
- 80% of the time: Replace the word with the [MASK] token, e.g., my dog is hairy → my dog is [MASK]
- 10% of the time: Replace the word with a random word, e.g., my dog is hairy → my dog is apple
- 10% of the time: Keep the word unchanged, e.g., my dog is hairy → my dog is hairy. The purpose of this is to bias the representation towards the actual observed word.

- Too little masking: too expensive to train
- Too much masking: not enough context

How is BERT pre-trained?

Two **pre-training** objectives:

- Masked language modeling (MLM)
- Next sentence prediction (NSP)

Sample two segments of text (**segment A** and **segment B**) and predict whether the second segment is followed after the first one.

Next sentence prediction

Input = [CLS] the man went to [MASK] store [SEP]
he bought a gallon [MASK] milk [SEP]

Label = IsNext

- 50% probability: a text segment of 512 tokens
- 50% probability: a text segment of 256 tokens, followed by another text segment of 256 tokens from a different document

Input = [CLS] the man [MASK] to the store [SEP]
penguin [MASK] are flight ##less birds [SEP]

Label = NotNext

The MLM loss and NSP loss are combined in pre-training

Remarks on BERT's pre-training objectives

- Later on, (Joshi et al., 2019; Liu et al., 2019) find that the next-sentence prediction objective unnecessary - RoBERTa doesn't use NSP at all.
- Understanding the role of masking rates (Why 15%)



Should You Mask 15% in Masked Language Modeling?

Alexander Wettig* Tianyu Gao* Zexuan Zhong Danqi Chen


Department of Computer Science, Princeton University

{awettig,tianyug,zzhong,danqic}@cs.princeton.edu

(Wetting et al., 2023)

- The optimal masking rate should depend on *model sizes* and *masking strategies*
- Masking plays two distinct roles: corruption vs prediction

How is BERT used for downstream tasks?



The pre-trained BERT encoder can be used directly for downstream tasks with **minimal task-specific parameters!**

BERT for text classification


"I love this movie.
I've seen it many times
and it's still awesome."



"This movie is bad.
I don't like it at all.
It's terrible."



- Add a classifier on top of the [CLS] representation
- New parameters: $d \times |C|$, jointly trained with BERT parameters
 - d = hidden dimension (e.g., 768)
 - $|C|$ = number of classes (e.g., 2)




BERT for question answering

Stanford Question Answering Dataset (SQuAD)

Tesla was the fourth of five children. He had an older brother named Dane and three sisters, Milka, Angelina and Marica. Dane was killed in a horse-riding accident when Nikola was five. In 1861, Tesla attended the "Lower" or "Primary" School in Smiljan where he studied German, arithmetic, and religion. In 1862, the Tesla family moved to Gospić, Austrian Empire, where Tesla's father worked as a pastor. Nikola completed "Lower" or "Primary" School, followed by the "Lower Real Gymnasium" or "Normal School."

Q: What language did Tesla study while in school?

A: German





$$\mathcal{L} = -\log p_{\text{start}}(s^*) - \log p_{\text{end}}(e^*)$$

$$p_{\text{start}}(i) = \text{softmax}_i(\mathbf{w}_{\text{start}}^\top \mathbf{h}_i)$$

$$p_{\text{end}}(i) = \text{softmax}_i(\mathbf{w}_{\text{end}}^\top \mathbf{h}_i)$$

Pre-BERT QA models vs BERT



All the parameters are pre-trained except for a small number of task-specific parameters $\mathbf{w}_{\text{start}}, \mathbf{w}_{\text{end}}$

Only the word embeddings at the input layer are pre-trained (2-3M parameters in total)


BERT has 110M or 330M parameters

Prompt-based fine-tuning

Task: sentiment classification

Input: “No reason to watch.”


Output: positive  or negative  ?



Fine-tuning

81.4% (32 examples)

93.5% (67k examples; BERT-base)



Prompt-based **fine-tuning**

92.7% (32 examples)

BERT: training cost

- BERT-base: 12 layers, $n = 512$, $d = 768$, 110M parameters
- BERT-large: 24 layers, $n = 512$, $d = 1024$, 330M parameters
- Trained on Wikipedia + BooksCorpus (3.3 billion tokens)
- Estimate: 6 days for BERT-base and 26 days for BERT-large on 8 Nvidia Titan-V GPUs (12Gb memory)

RESEARCH

by Jacob Portes*, Alex Trott*, Daniel King, Sam Havens
on March 9, 2023

SHARE



MosaicBERT: Pretraining BERT from Scratch for \$20

<https://www.mosaicml.com/blog/mosaicbert>

(8x A100-80Gb GPUs)

RoBERTa


- BERT is still under-trained
- Removed the next-sentence prediction objective
- Trained longer with 10x data & bigger batch sizes
- Pre-trained on 1,024 V100 GPUs for one day in 2019

| Model | data | bsz | steps | SQuAD (v1.1/2.0) | MNLI-m | SST-2 |
|--------------------------|-------|-----|-------|---------------------|-------------|-------------|
| RoBERTa | | | | | | |
| with BOOKS + WIKI | 16GB | 8K | 100K | 93.6/87.3 | 89.0 | 95.3 |
| + additional data (§3.2) | 160GB | 8K | 100K | 94.0/87.7 | 89.3 | 95.6 |
| + pretrain longer | 160GB | 8K | 300K | 94.4/88.7 | 90.0 | 96.1 |
| + pretrain even longer | 160GB | 8K | 500K | 94.6/89.4 | 90.2 | 96.4 |
| BERT _{LARGE} | | | | | | |
| with BOOKS + WIKI | 13GB | 256 | 1M | 90.9/81.8 | 86.6 | 93.7 |




ELECTRA

ELECTRA provides a more **efficient** training method, because it predicts 100% of tokens (instead of 15%) every time



BERT vs GPT-I models

- Unlike GPT-1 (autoregressive language models), BERT/RoBERTa/ELECTRA can't generate text naturally!
- However, bidirectionality is important for natural language understanding tasks.




(Devlin et al., 2019)

Why not combine the best
of both worlds?

T5:Text-to-text models

- T5 = **T**ext-to-**T**ext **T**ransfer **T**ransformer
- Transformer encoder-decoder architecture
 - Encoder: preserves bidirectionality
 - Decoder: good for generation!




T5:Text-to-text models

- T5 = Text-to-Text Transfer Transformer

Fine-tuning

All NLU tasks can be cast as text prediction tasks too!



T5:Text-to-text models

- T5 = **T**ext-**t**o-**T**ext **T**ransfer **T**ransformer

T5 comes in different sizes:

- t5-small.
- t5-base.
- t5-large.
- t5-3b.
- t5-11b.


Training corpus:

C4 = **C**olossal **C**lean **C**rawled **C**orpus

A cleaned version of Common Crawl

Trained on **34B** tokens

BART: another text-to-text model





The GPT-3 era: prompting and in-context learning

GPT-3

Language Models are Few-Shot Learners

Tom B. Brown* Benjamin Mann* Nick Ryder* Melanie Subbiah*

Jared Kaplan[†] Prafulla Dhariwal Arvind Neelakantan Pranav Shyam

Girish Sastry Amanda Askell Sandhini Agarwal Ariel Herbert-Voss

Gretchen Krueger Tom Henighan Rewon Child Aditya Ramesh

Daniel M. Ziegler Jeffrey Wu Clemens Winter

Christopher Hesse Mark Chen Eric Sigler Mateusz Litwin Scott Gray


Benjamin Chess Jack Clark Christopher Berner

Sam McCandlish Alec Radford Ilya Sutskever Dario Amodei

(Brown et al., 2020)

- From GPT-1 to GPT-2 to GPT-3
- **Transformer decoder** for autoregressive language modeling
- What's new?
 - Parameters ↑ Data ↑
 - These models can start generating human-like text!
 - You don't have to fine-tune these models - **prompting** works well!

GPT-1 vs GPT-2 vs GPT-3



<https://medium.com/@YanAlx/>

(Zha et al., 2023)

GPT-2 generation example

PROMPT
(HUMAN-WRITTEN)

A train carriage containing controlled nuclear materials was stolen in Cincinnati today. Its whereabouts are unknown.

MODEL COMPLETION

The incident occurred on the downtown train line, which runs from Covington and Ashland stations.

In an email to Ohio news outlets, the U.S. Department of Energy said it is working with the Federal Railroad Administration to find the thief.

“The theft of this nuclear material will have significant negative consequences on public and environmental health, our workforce and the economy of our nation,” said Tom Hicks, the U.S. Energy Secretary, in a statement. “Our top priority is to secure the theft and ensure it doesn’t happen again.”

The stolen material was taken from the University of Cincinnati’s Research Triangle Park nuclear research site, according to a news release from Department officials.


GPT-3's in-context learning

- Before GPT-3, most use of pre-trained language models are through ***fine-tuning***, on a reasonably-sized supervised dataset.
 - SST-2 has 67k examples, SQuAD has 88k (passage, answer, question) triples.
- GPT-3 shows that, with a very large **autoregressive language model (175B parameters)**, the model can perform a task:
 - Using only **a few examples** “**Few-shot learning**”
 - **Without gradient updates:**
the examples are only provided in the context! “**In-context learning**”


In-context learning

Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.



In-context learning



(Brown et al., 2020): Language Models are Few-Shot Learners


In-context learning

Input: 2014-06-01
Output: !06!01!2014!
Input: 2007-12-13
Output: !12!13!2007!
Input: 2010-09-23
Output: !09!23!2010!
Input: 2005-07-23
Output: !07!23!2005!

I
— model completion

in-context examples

test example



Understanding in-context learning

- **Hypothesis #1:** Transformers perform implicit gradient descent to update an “inner model”

Transformers Learn In-Context by Gradient Descent

Johannes von Oswald^{1,2} Egvind Niklasson² Ettore Randazzo² João Sacramento¹
Alexander Mordvintsev² Andrey Zhmoginov² Max Vladymyrov²


Why Can GPT Learn In-Context? Language Models Implicitly Perform Gradient Descent as Meta-Optimizers

Damai Dai^{†*}, Yutao Sun^{||*}, Li Dong[‡], Yaru Hao[‡], Shuming Ma[‡], Zhifang Sui[†], Furu Wei[‡]

- **Hypothesis #2:** Transformers learn tasks required for downstream applications during pre-training, and in-context demonstrations are only used to recognize which task is required

Rethinking the Role of Demonstrations: What Makes In-Context Learning Work?

Sewon Min^{1,2} Xinxix Lyu¹ Ari Holtzman¹ Mikel Artetxe²
Mike Lewis² Hannaneh Hajishirzi^{1,3} Luke Zettlemoyer^{1,2}
¹University of Washington ²Meta AI ³Allen Institute for AI
{sewon, alrope, ahai, hannaneh, lsz}@cs.washington.edu
{artetxe, mikelewis}@meta.com



Ground-truth labels don't matter!


Understanding in-context learning

We disentangle In-context learning into two roles -
task recognition (TR) vs task learning (TL)

- TR: recognizes the task from demonstrations and applies LLMs' pre-trained priors
- TL: learns a new input-label mapping from demonstrations
- ICL performs both TR and TL, but TL emerges with **larger models and more demonstrations**



Chain-of-thought (CoT) prompting






ChatGPT

The ChatGPT era:
Supervised instruction tuning and RLHF

From 2020's GPT-3 to 2022's ChatGPT



What's new?

- Code training
- Supervised instruction tuning
- RLHF = Reinforcement learning from human feedback

Post-training pipeline of LLMs

**Training language models to follow instructions
with human feedback**

Long Ouyang* **Jeff Wu*** **Xu Jiang*** **Diogo Almeida*** **Carroll L. Wainwright***

Pamela Mishkin* **Chong Zhang** **Sandhini Agarwal** **Katarina Slama** **Alex Ray**

John Schulman **Jacob Hilton** **Fraser Kelton** **Luke Miller** **Maddie Simens**

Amanda Askell†

Peter Welinder

Paul Christiano*†


Jan Leike*

Ryan Lowe*

OpenAI

The InstructGPT paper

Post-training pipeline of LLMs



Supervised fine-tuning (SFT)

- **Instruction data:** (instruction, response)
- **Objective:** next-token prediction conditioned on instruction

Can you write a short introduction about the relevance of the term "monopsony" in economics? Please use examples related to potential monopsonies in the labour market and cite relevant research.

"Monopsony" refers to a market structure where there is only one buyer for a particular good or service. In economics, this term is particularly relevant in the labor market, where a monopsony employer has significant power over the wages and working conditions of their employees. The presence of a monopsony can result in lower wages and reduced employment opportunities for workers, as the employer has little incentive to increase wages or provide better working conditions understand the extent and impact of monopsonies on the economy and to develop policies to address this issue. References: Bivens, J., & Mishel, L. (2013). The Pay of Corporate Executives and Financial Professionals as Evidence of Rents in Top 1 Percent Incomes. Journal of Economic Perspectives, 27(3), 57-78.


Supervised fine-tuning (SFT)

| Use-case | Prompt | Use-case | (%) |
|---------------|---|----------------|-------|
| Brainstorming | List five ideas for how to regain enthusiasm for my career | Generation | 45.6% |
| Generation | Write a short story where a bear goes to the beach, makes friends with a seal, and then returns home. | Open QA | 12.4% |
| Rewrite | This is the summary of a Broadway play: """ {summary} """ This is the outline of the commercial for that play: """ | Brainstorming | 11.2% |
| | | Chat | 8.4% |
| | | Rewrite | 6.6% |
| | | Summarization | 4.2% |
| | | Classification | 3.5% |
| | | Other | 3.5% |
| | | Closed QA | 2.6% |
| | | Extract | 1.9% |

SFT data: only ~13k (written by labeller, not public)

Collection of supervised fine-tuning data

- Repurposed from existing supervised datasets
 - Human-written instructions and responses




OpenAssisstant (Kopf et al., 2023)

Super-NaturalInstructions (Wang et al., 2022)


Also: FLAN, T0++

Collection of supervised fine-tuning data

- Response distilled from GPT models
- Instructions can be generated by GPT models too, e.g., **Self-Instruct** (Wang et al., 2023)



<https://sharegpt.com/>



Stanford Alpaca

Learning from human feedback


- **Preference data:** (instruction, winning response, losing response)

Preference Dataset: Signals for human desiderata



InstructGPT: 33k prompts, each with K (4~9) corresponding SFT model completions ranked by labellers

Reinforcement learning from human feedback




Multiple stages in training


- Reward model training
- Sampling from policy model
- Policy model update

Multiple models involved

- Reward model
- Policy model
- Reference model




GPT-3 vs InstructGPT



<https://openai.com/index/instruction-following/>

Direct preference optimization (DPO)

Instead of training an explicit reward model, express reward in the form of policy model:



Implicit reward expression:

$$r(x, y) = \beta \log \frac{\pi_\theta(y | x)}{\pi_{\text{ref}}(y | x)} + \beta \log Z(x)$$

Bradley-Terry ranking objective:

$$\mathcal{L}_R(r_\phi, \mathcal{D}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r_\phi(x, y_w) - r_\phi(x, y_l))]$$

DPO objective:

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_\theta(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \right]$$


Simple preference optimization (SimPO)

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_\theta(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \right]$$

$$\mathcal{L}_{\text{SimPO}}(\pi_\theta) = -\mathbb{E} \left[\log \sigma \left(\frac{\beta}{|y_w|} \log \pi_\theta(y_w | x) - \frac{\beta}{|y_l|} \log \pi_\theta(y_l | x) - \gamma \right) \right]$$

- A simple **length-normalized** reward (reference-free!):

$$r_{\text{SimPO}}(x, y) = \frac{\beta}{|y|} \log \pi_\theta(y | x) = \frac{\beta}{|y|} \sum_{i=1}^{|y|} \log \pi_\theta(y_i | x, y_{<i})$$



- Introducing target reward margin in Bradley-Terry objective:

$$p(y_w \succ y_l | x) = \sigma(r(x, y_w) - r(x, y_l) - \gamma)$$

Simple preference optimization (SimPO)

